



**Question 1:** Answer the following questions.

- Explain explicit, implicit, and automatic multi-threading, then write the pros and cons for each of them.
- Describe false sharing. Can you think of a change to the CPU architecture that would fix this issue? would it improve general performance of the CPU? (explain)
- Explain First Touch Policy in NUMA architectures and the way it influences parallel programs.

**Question 2:** The following program has been written by an LLM trying to calculate  $\int_{-5}^5 e^{-x^2} dx$  using OpenMP library. But we're getting inconsistent and wrong results. Fix and improve the performance of the given code and explain the problems in code.

```
#include <iostream>
#include <cmath>
#include <omp.h>

#define NUM_THREADS 8
#define NUM_STEPS 100000

int main() {
    static double step_len = 10.0 / (double) NUM_STEPS;
    double sum[NUM_THREADS] = {0.0};
    double result = 0.0;
    int num_threads;
    #pragma omp parallel num_threads(NUM_THREADS)
    {
        double x;
        int thread_id = omp_get_thread_num();
        num_threads = omp_get_num_threads();
        for (int i = thread_id; i < NUM_STEPS; i += NUM_THREADS) {
            x = (i + 0.5) * step_len - 5.0;
            sum[thread_id] += exp(-(x*x));
        }
    }
    for (int i = 0; i < num_threads; i++)
        result += sum[i] * step_len;
    std::cout << "Result: " << result << std::endl;
    return 0;
}
```

a)

b)

c) C

d)

e) (This section has double the points of other sections in this question)

```
int N;  
int mat[N][N];  
for (int i = 0; i < N; i++)  
    for (int j = 0; j < N; j++)  
#pragma omp parallel for  
        for (int k = 0; k < N; k++)  
#pragma omp critical  
            mat[i][j] += mat[(i+k)%N][(j+k)%N];
```

**Question 4:** The following function convolves the input 2D matrix with a 3x3 kernel and performs 2x2 pooling without using any extra dynamic (based on matrix size) memory allocation. For simplicity assume that matrix size is always even.

```
void conv_pool(int** mat, int** kernel, int N){
    int sum;
    for (int i = 1; i < N-1; i++){
        for (int j = 1; j < N-1; j++) {
            sum = 0;
            for (int k = -1; k < 2; k++)
                for (int l = -1; l < 2; l++)
                    sum += mat[i+k][j+l]
            mat[i-1][j-1] = sum;
        }
    }

    for (int i = 0; i < N-2; i+=2){
        for (int j = 0; j < N-2; j+=2) {
            a[i/2][j/2] = (a[i][j] + a[i+1][j] + a[i][j+1] + a[i+1][j+1]) / 4;
        }
    }

    // equivalent operation: mat = mat[0:(N-2)/2][0:(N-2)/2]
    resize_mat(mat, 0, (N-2)/2, 0, (N-2)/2);

    return;
}
```

- a) Using OpenMP parallelize this function and compare total memory usage in the serial and parallel versions. To simplify decomposition, assume that the number of threads is a square number and matrix size is an even multiple of square root of thread count.
- b) **(Bonus Question)** Is it possible to parallelize this function without use of any extra dynamic memory allocation? If yes, Implement the code otherwise prove why it cannot be done.